

Tidy Data

5 Most Common Problems With Messy Datasets

Pengfei Qin
12 April 2021

Why we need data tidying?

Data

- We use data for data analysis (regression or prediction), or do some larger projects (analyze pandemic trends).

Problems

- Data may have many disadvantages (missing values, outliers, typo).

color	director_name	duration	gross	movie_title	language	country	budget	title_year	imdb_score
Color	Martin Scorsese	240	116866727	The Wolf of Wall Street	English	USA	100000000	2013	8.2
Color	Shane Black	195	408992272	Iron Man 3	English	USA	200000000	2013	7.2
color	Quentin Tarantino	187	54116191	The Hateful Eight	English	USA	44000000	2015	7.9
Color	Kenneth Lonergan	186	46495	Margaret	English	usa	14000000	2011	6.5
Color	Peter Jackson	186	258355354	The Hobbit: The Desolation of Smaug	English	USA	225000000	2013	7.9
	N/A	183	330249062	Batman v Superman: Dawn of Justice	English	USA	250000000	2021	6.9
Color	Peter Jackson	-50	303001229	The Hobbit: An Unexpected Journey	English	USA	180000000	2012	7.9
Color	Edward Hall	180		Restless	English	UK		2012	7.2
Color	Joss Whedon	173	623279547	The Avengers	English	USA	220000000	2012	8.1
Color	Joss Whedon	173	623279547	The Avengers	English	USA	220000000	2012	8.1
	Tom Tykwer	172	27098580	Cloud Atlas	English	Germany	102000000	2012	-7.5
	Null	158	102515793	The Girl with the Dragon Tattoo	English	USA	90000000	2011	7.8
Color	Christopher Spencer	170	59696176	Son of God	English	USA	22000000	2014	5.6
Color	Peter Jackson	164	255108370	The Hobbit: The Battle of the Five Armies	English	New Zealand	250000000	2014	7.5
Color	Tom Hooper	158	148775460	Les Misérables	English	USA	61000000	2012	7.6
Color	Tom Hooper	158	148775460	Les Misérables	English	USA	61000000	2012	7.6

- The structure of the data set is not uniform (various arrangements in variables and observations).

Why we need data tidying?

Solution

- Data Tidying: Structuring datasets to facilitate analysis.
- Tidy data provides us a standard method to organize the values in a dataset.

Purpose

- Tidying the initial data can make subsequent data analysis easier.
- Consistent data structure will also make our data analysis easier.
- Make R happier.

The principle of tidy data

Tidy data

1. Every column is a variable.
2. Every row is an observation.
3. Each type of observational unit forms a table.

```
dataset4 = data.frame(Song = rep(c("Song_A", "Song_B", "Song_C"),  
  'Year' = rep(c("2019", "2020", "2021"), each = 3),  
  'Time' = rep(c("241", "175", "239"), each = 3),  
  'Date' = c("2021-04-01", "2021-04-02", "2021-04-03", "2021-04-01",  
  'Rank' = c(1,3,2,2,1,3,3,2,1))  
dataset4
```

##	Song	Year	Time	Date	Rank
## 1	Song_A	2019	241	2021-04-01	1
## 2	Song_A	2019	241	2021-04-02	3
## 3	Song_A	2019	241	2021-04-03	2
## 4	Song_B	2020	175	2021-04-01	2
## 5	Song_B	2020	175	2021-04-02	1
## 6	Song_B	2020	175	2021-04-03	3
## 7	Song_C	2021	239	2021-04-01	3
## 8	Song_C	2021	239	2021-04-02	2
## 9	Song_C	2021	239	2021-04-03	1

"5 Most Common Problems With Messy Datasets"

1. Column headers are values, not variable names.

```
dataset1 = data.frame(name = c("Company_A", "Company_B", "Company_C", "Company_D"),  
  'bachelor' = sample(1:10, 4),  
  'master' = sample(1:10, 4),  
  'PhD' = sample(1:10, 4))  
dataset1
```

```
##      name bachelor master PhD  
## 1 Company_A      4      4   5  
## 2 Company_B      7      9   4  
## 3 Company_C      5      5   1  
## 4 Company_D     10      7   9
```

1. Column headers are values, not variable names.

```
dataset1 <- dataset1 %>%  
  pivot_longer(-name, names_to = "degree", values_to = "frequency")  
dataset1
```

```
## # A tibble: 12 x 3  
##   name      degree frequency  
##   <chr>     <chr>      <int>  
## 1 Company_A bachelor      4  
## 2 Company_A master       4  
## 3 Company_A PhD          5  
## 4 Company_B bachelor      7  
## 5 Company_B master       9  
## 6 Company_B PhD          4  
## 7 Company_C bachelor      5  
## 8 Company_C master       5  
## 9 Company_C PhD          1  
## 10 Company_D bachelor     10  
## 11 Company_D master       7  
## 12 Company_D PhD          9
```

This form is tidy: there's one variable in each column, and each row represents one observation.

2. Multiple variables stored in one column

```
dataset2 = data.frame(name = c("Company_A", "Company_B", "Compar  
'm2035' = sample(1:10, 4),  
'm3550' = sample(1:10, 4),  
'm5065' = sample(1:10, 4),  
'f2035' = sample(1:10, 4),  
'f3550' = sample(1:10, 4),  
'f5060' = sample(1:10, 4))  
dataset2
```

```
##      name m2035 m3550 m5065 f2035 f3550 f5060  
## 1 Company_A     8    10     1     8     8     8  
## 2 Company_B    10     4     8     7     6     2  
## 3 Company_C     3     9     5     3     9     5  
## 4 Company_D     6     2     3     2     1     7
```


2. Multiple variables stored in one column

```
dataset2 <- dataset2 %>%  
  pivot_longer(-name, names_to = "combination", values_to = "frequency")  
dataset2
```

```
## # A tibble: 24 x 3  
##   name      combination frequency  
##   <chr>      <chr>          <int>  
## 1 Company_A m2035             8  
## 2 Company_A m3550            10  
## 3 Company_A m5065             1  
## 4 Company_A f2035             8  
## 5 Company_A f3550             8  
## 6 Company_A f5060             8  
## 7 Company_B m2035            10  
## 8 Company_B m3550             4  
## 9 Company_B m5065             8  
## 10 Company_B f2035             7  
## # ... with 14 more rows
```

2. Multiple variables stored in one column

```
dataset2 <- dataset2 %>%  
  separate(combination, c("sex", "age"),1)  
dataset2
```

```
## # A tibble: 24 x 4  
##   name      sex  age frequency  
##   <chr>    <chr> <chr>    <int>  
## 1 Company_A m    2035      8  
## 2 Company_A m    3550     10  
## 3 Company_A m    5065      1  
## 4 Company_A f    2035      8  
## 5 Company_A f    3550      8  
## 6 Company_A f    5060      8  
## 7 Company_B m    2035     10  
## 8 Company_B m    3550      4  
## 9 Company_B m    5065      8  
## 10 Company_B f    2035      7  
## # ... with 14 more rows
```

This form is tidy: there's one variable in each column, and each row represents one observation.

3. Variables are stored in both columns and rows.

```
dataset3 = data.frame(city = rep(c("Beijing", "Hong Kong", "Los
  'month'= c("January"),
  'element'= c("avg_environmental_quality", "avg_air_quality"),
  'y2019'= sample(c("high", "median", "low"),8, replace = TRUE),
  'y2020'= sample(c("high", "median", "low"),8, replace = TRUE),
  'y2021'= sample(c("high", "median", "low"),8, replace = TRUE))
dataset3
```

##	city	month	element	y2019	y2020	y2021
## 1	Beijing	January	avg_environmental_quality	high	low	high
## 2	Beijing	January	avg_air_quality	high	high	high
## 3	Hong Kong	January	avg_environmental_quality	high	median	low
## 4	Hong Kong	January	avg_air_quality	median	low	high
## 5	Los Angeles	January	avg_environmental_quality	median	high	high
## 6	Los Angeles	January	avg_air_quality	high	median	median
## 7	New York	January	avg_environmental_quality	high	low	high
## 8	New York	January	avg_air_quality	high	low	low

3. Variables are stored in both columns and rows.

```
dataset3 <- dataset3 %>%
  pivot_longer(y2019:y2021, names_to = "year", values_to = "value")

dataset3 <- dataset3 %>%
  mutate(year = as.integer(gsub("y", "", year))) %>%
  select(city, month, element, year, value)

dataset3
```

```
## # A tibble: 24 x 5
##   city      month  element      year value
##   <chr>    <chr>   <chr>      <int> <chr>
## 1 Beijing January avg_environmental_quality 2019 high
## 2 Beijing January avg_environmental_quality 2020 low
## 3 Beijing January avg_environmental_quality 2021 high
## 4 Beijing January avg_air_quality      2019 high
## 5 Beijing January avg_air_quality      2020 high
## 6 Beijing January avg_air_quality      2021 high
## 7 Hong Kong January avg_environmental_quality 2019 high
## 8 Hong Kong January avg_environmental_quality 2020 median
## 9 Hong Kong January avg_environmental_quality 2021 low
## 10 Hong Kong January avg_air_quality      2019 median
## # ... with 14 more rows
```

3. Variables are stored in both columns and rows.

```
dataset3 <- dataset3 %>%
  pivot_wider(names_from = element, values_from = value)
dataset3
```

```
## # A tibble: 12 x 5
##   city      month    year avg_environmental_quality avg_air_qua
##   <chr>    <chr>  <int> <chr>                        <chr>
## 1 Beijing January  2019 high                       high
## 2 Beijing January  2020 low                        high
## 3 Beijing January  2021 high                       high
## 4 Hong Kong January  2019 high                       median
## 5 Hong Kong January  2020 median                     low
## 6 Hong Kong January  2021 low                        high
## 7 Los Angeles January  2019 median                     high
## 8 Los Angeles January  2020 high                       median
## 9 Los Angeles January  2021 high                       median
## 10 New York January  2019 high                       high
## 11 New York January  2020 low                        low
## 12 New York January  2021 high                       low
```

This form is tidy: there's one variable in each column, and each row represents one observation.

4. Multiple types of observational units are stored in the same table

##		Song	Year	Time	Date	Rank
##	1	Song_A	2019	241	2021-04-01	1
##	2	Song_A	2019	241	2021-04-02	3
##	3	Song_A	2019	241	2021-04-03	2
##	4	Song_B	2020	175	2021-04-01	2
##	5	Song_B	2020	175	2021-04-02	1
##	6	Song_B	2020	175	2021-04-03	3
##	7	Song_C	2021	239	2021-04-01	3
##	8	Song_C	2021	239	2021-04-02	2
##	9	Song_C	2021	239	2021-04-03	1

```
song <- dataset4 %>%
  distinct(Song, Year, Time)
song
```

##		Song	Year	Time
##	1	Song_A	2019	241
##	2	Song_B	2020	175
##	3	Song_C	2021	239

4. Multiple types of observational units are stored in the same table

```
rank <- dataset4 %>%  
  left_join(song, c("Song", "Year", "Time")) %>%  
  select(Song, Date, Rank)  
rank
```

##		Song	Date	Rank
##	1	Song_A	2021-04-01	1
##	2	Song_A	2021-04-02	3
##	3	Song_A	2021-04-03	2
##	4	Song_B	2021-04-01	2
##	5	Song_B	2021-04-02	1
##	6	Song_B	2021-04-03	3
##	7	Song_C	2021-04-01	3
##	8	Song_C	2021-04-02	2
##	9	Song_C	2021-04-03	1

5. A single observational unit is stored in multiple tables

```
GDP_and_Tax = data.frame(City = rep(c("Beijing", "Hong Kong", "New York"), 2),
  'GDP' = runif(6, 100, 200),
  'Tax_Revenue' = runif(6, 15, 25))
GDP_and_Tax
```

```
##           City      GDP Tax_Revenue
## 1   Beijing 180.6807    22.73900
## 2   Beijing 181.1368    21.24779
## 3 Hong Kong 108.6159    21.25485
## 4 Hong Kong 156.1439    18.90633
## 5 New York 111.7318    19.66630
## 6 New York 112.4898    24.78881
```

```
Energy_and_Industry = data.frame(City = rep(c("Beijing", "Hong Kong", "New York"), 2),
  'Energy_Consumption' = runif(6, 1900, 2000),
  'Industrial_Output' = runif(6, 250, 300))
Energy_and_Industry
```

```
##           City Energy_Consumption Industrial_Output
## 1   Beijing      1909.612           265.5763
## 2   Beijing      1913.279           280.9248
## 3 Hong Kong      1925.124           279.3627
## 4 Hong Kong      1932.035           263.0203
## 5 New York       1985.700           286.1806
## 6 New York       1942.070           285.3541
```


5. A single observational unit is stored in multiple tables

```
dataset5 <- inner_join(GDP_and_Tax, Energy_and_Industry)
```

```
## Joining, by = "City"
```

```
dataset5
```

##	City	GDP	Tax_Revenue	Energy_Consumption	Industrial_Output
## 1	Beijing	180.6807	22.73900	1909.612	265.
## 2	Beijing	180.6807	22.73900	1913.279	280.
## 3	Beijing	181.1368	21.24779	1909.612	265.
## 4	Beijing	181.1368	21.24779	1913.279	280.
## 5	Hong Kong	108.6159	21.25485	1925.124	279.
## 6	Hong Kong	108.6159	21.25485	1932.035	263.
## 7	Hong Kong	156.1439	18.90633	1925.124	279.
## 8	Hong Kong	156.1439	18.90633	1932.035	263.
## 9	New York	111.7318	19.66630	1985.700	286.
## 10	New York	111.7318	19.66630	1942.070	285.
## 11	New York	112.4898	24.78881	1985.700	286.
## 12	New York	112.4898	24.78881	1942.070	285.